

COMP 3500 Introduction to Operating Systems

Project 5 - CPU Scheduling

Long Version: 1.4
4/27/2023

Points Possible: 100
Submission via Canvas

This is an individual assignment; no collaboration among students.

1. Learning Objectives

The goal of this project is to implement a simple CPU schedule, which orchestrates three CPU scheduling policies. You will achieve the following objectives upon the completion of the project.

- To design a simple CPU scheduler
- To implement three scheduling policies in C/C++
- To conduct scheduling simulation experiments on a Linux machine
- Use GDB to debug your C/C++ program

2. Requirements

You will implement a CPU scheduler that selects a job according to a given scheduling policy. Because the project intends to simulate a CPU scheduler, there is no need for any actual process creation, execution, or termination. When a task is scheduled, the simulator simply prints out what task is picked to run at a time. The outputs of your scheduler are similar to the Gantt chart style.

- Each student should independently accomplish this project assignment. You may discuss with other students to solve the coding problems. Students should not share any project code with any other student. Collaborations among students in any form will be treated as a serious violation of the University's academic integrity code.
- **CAUTION!** You must follow the specified output format.
- You must write a C/C++ program on the Linux operating system.

3. Implementing a CPU Scheduler

3.1 Scheduling Policies

Your simulator should manage the following three scheduling policies. The detailed algorithms will be introduced and discussed in COMP 3500 lectures.

- First Come First Serve (FCFS),
- Round Robin or RR, and
- Shortest Remaining Time First or SRTF).

3.2 Task Information and an Input File

The task information will be read from an input text file, the format of which is

```
pid arrival_time burst_time
```

All of fields are integer type where

- `pid` is a unique numeric process ID
- `arrival_time` is the time instance at which a task arrives
- `burst_time` is the CPU time requested by a task

The time unit of `arrival_time`, `burst_time`, and `interval` is millisecond.

3.3 Command-line Usage

Users should run your scheduler in a command-line as:

```
$scheduler task_list_file [FCFS|RR|SRTF] [time_quantum]
```

where `input_file` is the file name with task information described in Section 3.2 on page 1. `FCFS`, `RR`, and `SRTF` are names of the scheduling policies (a.k.a., algorithms). Parameter `time_quantum` is only applicable to the `RR` policy. The `FCFS` policy is nonpreemptive whereas `RR` and `SRTF` are all preemptive. Table below summarize three example usages.

Sample Usage	Description
<code>scheduler list1 FCFS</code>	Simulate the FCFS policy with a task file named “input1”
<code>scheduler list2 RR 5</code>	Simulate RR with the “input2” task file; time quantum is set to 5.
<code>scheduler list3 SRTF</code>	Simulate the SRTF policy with a task file named “input3”

3.4 System Design

Data Flow Diagram. You are suggested to start the design of your scheduler by crafting a data flow diagram, which enables you to decide what modules should be implemented and what are correlations among these modules. With the data flow diagram in place, you are expected to have a clear picture on key data structures used in your scheduler. In the data-flow-diagram design phase, please ignore any issue related to algorithm design (e.g., time-driven simulation, scheduling policies).

Loading Input File. After the simulator reads a task information list from an input text file, all the task information should be stored in a task list (e.g., an array, a dynamic array, or a linked list). Then, the simulator schedules tasks according to a scheduling policy specified in the command-line.

Time-Driven Simulation. Your simulator manages tasks in a time-driven manner. More specifically, at each time unit (a.k.a., time slot), your simulator inserts any newly arrived task into the ready queue, followed by invoking the scheduler to choose an appropriate task from the ready queue. When a task is elected to run, the simulator prints out a message indicating the process ID of a task running in this time slot. If no task is running (i.e., the ready queue is empty) during the current time unit, the simulator prints out an “idle” message. Before advancing to the next time unit, the simulator should update all necessary changes in the running task and the ready queue status.

3.5 Output and Statistical Information

Your simulator will be applied to compare the performance among the three scheduling policies. The scheduler is expected to compute the following statistical information referred to performance metrics. You will learn these metrics in our lectures. Please refer to Section 6 on page 4 for a sample output.

- average waiting time
- average response time
- average turnaround time
- overall CPU usage

3.6 Suggested Functions

You are recommended to implement the following functions.

- **A command-line parser.** Please refer to the command-line parser in os161 as an example.
- **Three scheduling algorithms.** Each scheduling policy should be implemented in a dedicated function.
- **Read an input file.** This function is responsible for loading a task information list from an input file into your simulator.
- **Print statistic information.** All the performance measures (e.g., average waiting time) must be displayed upon the completion of a simulation.
- **Error handler.** If an input command doesn't follow the specified format, the error handler must handle all types of input errors.

4. Programming Requirements

4.1 Programming Environment

Write your simulated virtual memory system in either C or C++. Compile and run your system using the gcc or g++ compiler on a Linux box (either in the computer labs in Shelby, your home Linux machine, a Linux box on a virtual machine, or using an emulator like Cygwin).

4.2 Function-Oriented Approach

You are *strongly suggested* to use a structure-oriented (a.k.a., function-oriented) approach for this project. In other words, you will need to write function definitions and use those functions; you can't just throw everything in the `main()` function. A well-done implementation will produce a number of robust functions, many of which may be useful for future programs in this project and beyond. Remember good design practices include:

- A function should do one thing, and do it well
- Functions should NOT be highly coupled

4.3 File Names

You will lose points if you do not use the specific program file name, or do not have a comment block on every source code file you hand in.

4.4 Usability Concerns and Error-Checking

You should appropriately prompt your user and assume that they only have basic knowledge of the system. You should provide enough error-checking that a moderately informed user will not crash your system. This should be discovered through your unit-testing. Your prompts should still inform the user of what is expected of them, even if you have error-checking in place.

5. Separate Compilation

You must use separate compilation and create a makefile for this project.

5.1 What's Make?

Make is a program that looks for a file called "makefile" or "Makefile", within the makefile are variables and things called dependencies. There are many things you can do with makefiles, if all you've ever done with makefiles is compile C or C++ then you are missing out. Pretty much anything that needs to be compiled (postscript, java, Fortran), can utilize makefiles.

5.2 Format of Makefiles -- Variables

First, let's talk about the simpler of the two ideas in makefiles, variables. Variable definitions are in the following format:

```
VARNAME = Value
```

So let's say I want to use a variable to set what compiler I'm going to use. This is helpful b/c you may want to switch from cc to gcc or to g++. We would have the following line in our makefile

```
CC = gcc
```

This assigns the variable CC to the string "gcc". To expand variables, use the following form:

```
${VARNAME}
```

So to expand our CC variable we would say:

```
${CC}
```

5.3 Format of Makefiles -- Dependencies

Dependencies are the heart of makefiles. Without them nothing would work. Dependencies have the following form:

```
dependency1: dependencyA dependencyB ... dependencyN  
command for dependency1
```

Check out the following links for more information on makefiles:

https://www.gnu.org/software/make/manual/html_node/Introduction.html

6. Sample Input and Output

Suppose an input task file is called "task.list", your scheduler should run as follows to simulate the three scheduling policies.

```
%more task.list
```

```
1 0 10  
2 0 9  
3 3 5  
4 7 4  
5 10 6  
6 10 7
```

```
%scheduler
```

```
Usage: scheduler task_list_file [FCFS|RR|SRTF] [time_quantum]
```

```
%scheduler task.list FCFS
```

```
Scheduling Policy: FCFS
```

```
There are 6 tasks loaded from "task.list". Press any key to continue ...
```

```
=====
```

```
<time 0> process 1 is running  
<time 1> process 1 is running  
<time 2> process 1 is running  
<time 3> process 1 is running  
<time 4> process 1 is running  
<time 5> process 1 is running  
<time 6> process 1 is running  
<time 7> process 1 is running  
<time 8> process 1 is running
```

```

<time 9> process 1 is running
<time 10> process 1 is finished...
<time 10> process 2 is running
<time 11> process 2 is running
<time 12> process 2 is running
<time 13> process 2 is running
<time 14> process 2 is running
<time 15> process 2 is running
<time 16> process 2 is running
<time 17> process 2 is running
<time 18> process 2 is running
<time 19> process 2 is finished...
<time 19> process 3 is running
<time 20> process 3 is running
<time 21> process 3 is running
<time 22> process 3 is running
<time 23> process 3 is running
<time 24> process 3 is finished...
<time 24> process 4 is running
<time 25> process 4 is running
<time 26> process 4 is running
<time 27> process 4 is running
<time 28> process 4 is finished...
<time 28> process 5 is running
<time 29> process 5 is running
<time 30> process 5 is running
<time 31> process 5 is running
<time 32> process 5 is running
<time 33> process 5 is running
<time 34> process 5 is finished...
<time 34> process 6 is running
<time 35> process 6 is running
<time 36> process 6 is running
<time 37> process 6 is running
<time 38> process 6 is running
<time 39> process 6 is running
<time 40> process 6 is running
<time 41> process 6 is finished...
<time 41> All processes finish .....

```

```

=====
Average waiting time:    14.17
Average response time:  14.17
Average turnaround time: 21.00
Overall CPU usage:      100.00%
=====

```

```
%scheduler
```

```
Usage: scheduler task_list_file [FCFS|RR|SRTF] [time_quantum]
```

```
%scheduler task.list SRTF
```

```
Scheduling Policy: SRTF
```

```
There are 6 tasks loaded from "task.list". Press any key to continue ...
```

```

=====
<time 0> process 2 is running
<time 1> process 2 is running

```

```

<time 2> process 2 is running
<time 3> process 3 is running
<time 4> process 3 is running
<time 5> process 3 is running
<time 6> process 3 is running
<time 7> process 3 is running
<time 8> process 3 is finished...
<time 8> process 4 is running
<time 9> process 4 is running
<time 10> process 4 is running
<time 11> process 4 is running
<time 12> process 4 is finished...
<time 12> process 2 is running
<time 13> process 2 is running
<time 14> process 2 is running
<time 15> process 2 is running
<time 16> process 2 is running
<time 17> process 2 is running
<time 18> process 2 is finished...
<time 18> process 5 is running
<time 19> process 5 is running
<time 20> process 5 is running
<time 21> process 5 is running
<time 22> process 5 is running
<time 23> process 5 is running
<time 24> process 5 is finished...
<time 24> process 6 is running
<time 25> process 6 is running
<time 26> process 6 is running
<time 27> process 6 is running
<time 28> process 6 is running
<time 29> process 6 is running
<time 30> process 6 is running
<time 31> process 6 is finished...
<time 31> process 1 is running
<time 32> process 1 is running
<time 33> process 1 is running
<time 34> process 1 is running
<time 35> process 1 is running
<time 36> process 1 is running
<time 37> process 1 is running
<time 38> process 1 is running
<time 39> process 1 is running
<time 40> process 1 is running
<time 41> process 1 is finished...
<time 41> All processes finish .....

```

```

=====
Average waiting time:  10.50
Average response time:  9.00

```

Average turnaround time:17.33

Overall CPU usage : 100.00 %

=====

%scheduler

Usage: scheduler task_list_file [FCFS|RR|SRTF] [time_quantum]

%scheduler task.list RR 4

Scheduling Policy: RR

There are 6 tasks loaded from "task.list". Press any key to continue ...

=====

<time 0> process 1 is running
<time 1> process 1 is running
<time 2> process 1 is running
<time 3> process 1 is running
<time 4> process 2 is running
<time 5> process 2 is running
<time 6> process 2 is running
<time 7> process 2 is running
<time 8> process 3 is running
<time 9> process 3 is running
<time 10> process 3 is running
<time 11> process 3 is running
<time 12> process 1 is running
<time 13> process 1 is running
<time 14> process 1 is running
<time 15> process 1 is running
<time 16> process 4 is running
<time 17> process 4 is running
<time 18> process 4 is running
<time 19> process 4 is running
<time 20> process 4 is finished...
<time 20> process 2 is running
<time 21> process 2 is running
<time 22> process 2 is running
<time 23> process 2 is running
<time 24> process 5 is running
<time 25> process 5 is running
<time 26> process 5 is running
<time 27> process 5 is running
<time 28> process 6 is running
<time 29> process 6 is running
<time 30> process 6 is running
<time 31> process 6 is running
<time 32> process 3 is running
<time 33> process 3 is finished...
<time 33> process 1 is running
<time 34> process 1 is running
<time 35> process 1 is finished...
<time 35> process 2 is running

```

<time 36> process 2 is finished...
<time 36> process 5 is running
<time 37> process 5 is running
<time 38> process 5 is finished...
<time 38> process 6 is running
<time 39> process 6 is running
<time 40> process 6 is running
<time 41> process 6 is finished...
<time 41> All processes finish .....
=====
Average waiting time:    22.00
Average response time:   8.33
Average turnaround time:28.83
Overall CPU usage:       100.00%
=====

```

7. Project Report

Write a project report that explains how you design and implement your virtual memory manager. Your report must address all of the questions. Your report is worth 15 points.

CAUTION! Your project report should provide sample input and output from your scheduler.

7.1 (6 points) Solution Description

You must describe your solutions with good sample input and output to show users what your scheduler is doing.

- (1) How did you separate scheduling mechanism from scheduling policies?
- (2) How did implement the three scheduling algorithms?
- (3) How did you calculate waiting times?
- (4) How did you calculate response times?
- (5) How did you calculate turnaround times?
- (6) How did you implement the command-line parser?

7.2 (6 points) Generality and Error Checking

- (1) How general is your solution?
- (2) How easy would it be to add a new scheduling policy into your scheduler?
- (3) Does your program offer input error checking?

7.3 (3 points) Miscellaneous factors

- (1) Is your code elegant or readable?
- (2) How innovative is your solution? Did you try any ideas not suggested here?
- (3) Did you document all outside sources?

8. Deliverables

8.1 Final Submission

Your final submission should include:

- A project report (see also Section 7)
- A copy of the complete source code – all source code files - of your CPU scheduler.

- Makefile (see also Section 5 on page 3)

CAUTION! You must submit a single compressed file (see Section 8.2) as a .tgz file, which includes both a project report, all source code files, and Makefile.

8.2 A Single Compressed File: `project5.tgz`

Now, submit your tarred and compressed file named `project5.tgz` through Canvas. You must submit your single compressed file through Canvas (no e-mail submission is accepted).

8.3 What happens if you can't complete the project?

If you are unable to complete this project for some reasons, please describe in your final design document the work that remains to be finished. It is important to present an honest assessment of any incomplete components.

9. Grading Criteria

The approximate marks allocation will be:

- 1) Project Report: 15%
- 2) Implementation (i.e., Source Code): 60%
- 3) Makefile: 15%
- 4) Clarity and attention to details (e.g., Adhering to coding and comment styles): 10%

10. Late Submission Penalty

- Ten percent (10%) penalty per day for late submission. For example, an assignment submitted after the deadline but up to 1 day (24 hours) late can achieve a maximum of 90% of points allocated for the assignment. An assignment submitted after the deadline but up to 2 days (48 hours) late can achieve a maximum of 80% of points allocated for the assignment.
- Assignment submitted more than 3 days (72 hours) after the deadline will not be graded.

11. Rebuttal period

- You will be given a period a week (7 days) to read and respond to the comments and grades of your homework or project assignment. The TA may use this opportunity to address any concern and question you have. The TA also may ask for additional information from you regarding your homework or project.